



RideMetrics: Optimizing Ride-Hailing Performance Across Uber & Bolt

Project Overview

RideMetrics is a personal data analytics project based on over one year of real-world data collected while driving for Bolt and Uber in Lodz, Poland. Motivated by the need to maximize earnings while balancing academic commitments. This project uses personal trip data to answer key business questions around earnings optimization and demand patterns. Passengers' location data has been excluded due to data protection reasons. The aim is to transform raw data into a clean, structured and actionable business insight using data cleaning, exploratory analysis, time series techniques, and an interactive dashboards.

Objectives & Business Questions

To optimize my earnings by identifying the most profitable working patterns and building a machine learning model to predict the best hours and days to work- transforming over one year of personal trip data into actionable Business Insights.

S/N	Business Questions	Datasets
1	What hours of the day generates the highest earnings?	Uber
2	Which day of the week is most profitable?	Uber
3	What is weekly total earnings?	Uber
4	What is the monthly trends for over 16months	Uber
5	Which month of the year is most profitable?	Uber
6	What is the average Uber deduction rate per trip?	Uber
7	What is average earnings per km?	Uber
8	How has monthly revenue and net earnings trended for over 16months?	Uber
9	Which hours have the highest trip demand?	Uber & Bolt
10	Which days have the highest trip volume?	Uber & Bolt
11	During which seasons is demand highest?	Uber & Bolt
12	Can historical patterns predict High/Medium/Low demand periods?	Uber & Bolt
13	Can historical patterns predict High/Medium/Low earning periods?	Uber

Data Description

Data Source

The dataset was personally collected through active driving on two ride-hailing platforms — **Uber** and **Bolt** — in Lodz, Poland. The data spans over one year of real-world trip activity and was exported directly from both platform apps in CSV format.

Data Fields

Property	Details
Format	CSVs
Platforms	Bolt & Uber
Location	Lodz Poland
Coverage	Over one year (Sep 2024 - Feb 2026)

Data Fields

Uber — Trip Details Table

- `global_product_name`
- `status`
- `request_timestamp_local`
- `begintrip_timestamp_local`
- `dropoff_timestamp_local`
- `trip_distance_miles`
- `trip_duration_seconds`
- `base_fare_local`
- `original_fare_local`
- `cancellation_fee_local`

Uber — Payment Table

- `City Name`
- `Trip UUID`
- `Local Amount`
- `Classification`
- `Category`
- `Local Timestamp`

Excluded Fields

- `vehicle_uuid` and `license_plate` removed from Trip Details table for privacy reasons.
- `Trip UUID` and `Currency Code` dropped from both tables as they add no analytical value.

Joining Key

Both Uber tables will be joined on:

- `request_time`
- `Local Timestamp`

In the Silver layer to create one unified Uber trips table.

Bolt Data Fields

Bolt — Trip Details Table

- `Time`
- `Driver Accepted`
- `Ride Finished`
- `Ride Distance`
- `Ride Duration`
- `Order State`

Bolt — Payment Table

- `Date`
- `Payment Method`
- `Date of Ride`
- `Price (no VAT)`
- `VAT`
- `Price Total`

Excluded Fields

- `Driver Name`, `Driver Phone`, `Driver Company`, `Vehicle Model` and `Car Reg Number` removed from Trip Details table for driver privacy reasons.

- `Driver App Version` , `Pickup Distance` , `Pickup Duration` , `Reject Reason` , `Estimated Pickup Time` , `Estimated Distance` and `Estimated Duration` dropped as they add no analytical value.
- `Country` dropped as all trips are in Poland.
- Bolt Payment table loaded into Bronze layer but subsequently excluded from Silver layer after investigation revealed incomplete and discounted payment data — retained for trip volume analysis only.

Data Limitations

- No location data limits geographic and route-level analysis as passenger pickup and drop-off coordinates were excluded for GDPR compliance.
- Cancelled trips with cancellation fees are retained in all revenue reporting tables as they represent legitimate earnings but excluded from ML feature tables as zero distance and duration provide no meaningful trip characteristic signal.
- Data reflects a single driver's experience in Łódź, Poland and may not be generalised to all drivers or other cities.
- External factors such as weather, traffic conditions and local events are not captured in the dataset and may influence demand and earnings patterns.
- Bolt trip data contains incomplete and discounted payment information and has been excluded from all financial analysis. Bolt data is used solely for time-based and operational pattern analysis.
- Bolt payment table was loaded into Bronze layer with 50,561 records but subsequently excluded from the Silver layer after investigation revealed three independent data quality issues — discounted customer payments only, inconsistent payment recording times by payment method and date of ride values not corresponding to actual trip dates.
- Uber driver incentives are included in net earnings used for ML model training as they form part of actual earnings received and historically correlate with specific time periods. This approach may overestimate earnings on days when Uber no longer runs incentive programmes.
- 12 special event days were identified as statistical outliers and flagged with a binary event indicator in the ML earnings features table. These dates

were verified against Łódź event records including concerts at Atlas Arena, cultural events and public holidays.

Data Quality Findings

Methodology

RideMetrics follows a structured analytics workflow built around the **Medallion Architecture** in Microsoft SQL Server, progressing data from raw ingestion through to business-ready insights.

Medallion Architecture (SQL Layer)

Layer	Purpose	Tables
1.  Bronze	Raw CSV data loaded as-is from Uber & Bolt	<code>Bronze.Uber_Trips</code> , <code>Bronze.Uber_Payments</code> , <code>Bronze.Bolt_Trips</code> , <code>Bronze.Bolt_Payments</code>
2.  Silver	Cleaned, standardized and merged trip data with engineered features	<code>Silver.Uber</code> , <code>Silver.Uber_Payments_Silver_pivoted</code> , <code>Silver.Bolt</code> , <code>Silver.Combined_Silver</code>
3.  Gold	Aggregated business-ready tables feeding Power BI and ML model	<code>Gold.Daily_Revenue</code> , <code>Gold.Hourly_Revenue</code> , <code>Gold.Weekly_Revenue</code> , <code>Gold.Monthly_Revenue</code> , <code>Gold.Yearly_Monthly_Revenue</code> , <code>Gold.Day_of_Week_Revenue</code> , <code>Gold.Hourly_Demand</code> , <code>Gold.Daily_Demand</code> , <code>Gold.Seasonal_Demand</code> , <code>Gold.ML_Demand</code> , <code>Gold.ML_Earnings</code>

Machine Learning

RideMetrics adopted two machine learning models, each addressing a distinct business question. Three algorithms were trained and compared for each model to ensure the best performing model was selected. XGBoost was selected as the best model for demand prediction achieving 95.7% accuracy while Random Forest was selected as the best model for earnings prediction achieving 76.7% accuracy.

Algorithm	Role	Demand Model	Earnings Model
Logistic Regression	Baseline model	47.1%	75.7%
Random Forest Classifier	Primary model	80.3%	76.7% 
XGBoost	Advanced benchmark	95.7% 	74.4%

Model 1: Demand Pattern Prediction

This model predicts whether a given hour and day combination represents a High, Medium, or Low demand period based on historical trip volume patterns. Both Uber and Bolt datasets were used since demand prediction relies on trip volume and time-based features rather than fare values. Demand categories were defined using a scoring system combining hour of day, day of week and season — reflecting patterns identified during EDA.

Target Variable: Demand category — High / Medium / Low

Input Features: Hour of day, day of week, month, season, platform, distance, duration

Use Case: *"Is it worth going online right now based on expected trip demand?"*

Best Algorithm: XGBoost

Results: Accuracy 95.7% — F1 Score 95.3% — CV Mean F1 92.3%

Model 2: Earnings Pattern Prediction

This model predicts whether a given hour represents a High, Medium, or Low earning period based on historical hourly net earnings patterns. Only Uber data was used since Bolt payment data is excluded from all financial analysis.

Earnings were aggregated at hourly level as the goal is to predict how much can be earned per hour rather than per trip. Thresholds were defined using the 33rd and 66th percentiles — Low below 26.14 PLN per hour, Medium between 26.14 and 41.81 PLN per hour and High above 41.81 PLN per hour.

Target Variable: Earnings category — High / Medium / Low

Input Features: Hour of day, day of week, month, season, average distance, average duration, total trips per hour, demand score

Use Case: *"Is it worth going online right now based on expected earnings?"*

Best Algorithm: Random Forest

Results: Accuracy 76.7% — F1 Score 76.4% — CV Mean F1 75.0% (Std 2.7%)

Combined Model Output

The true value of both models lies in combining their predictions into a single shift recommendation. By running both models simultaneously a driver can make a fully informed decision about whether to go online at any given hour:

-  **High Demand + High Earnings** → Definitely work — peak period with maximum returns
-  **High Demand + Low Earnings** → Work for trip volume — busy period but lower individual fares
-  **Low Demand + High Earnings** → Work for quality trips — fewer trips but higher value per trip
-  **Low Demand + Low Earnings** → Stay home — not worth going online

Analysis Phases

Phase	Description	Tools
1. Data Ingestion	Load Raw CSVs into Bronze layers	Python(Pandas,pyodbc),MSSQL
2. Data Cleaning and Feature Engineering	Standardize, merge, and enrich data in Silver layer	Python (pandas), MSSQL
3. EDA	Uncover patterns in earnings, demand, and trip volume across time periods, days, months and seasons	Python(Matplotlib, Seaborn)
4. Machine Learning	XGBoost Classifier selected for Demand Prediction (Uber + Bolt) achieving 95.7% accuracy. Random Forest Classifier selected for Earnings Prediction (Uber only) achieving 76.7% accuracy. Logistic Regression used as baseline for both models.	Python (Scikit-Learn, XGBoost)
5. Dashboard	Interactive visualizations connected directly to Gold layer tables in MSSQL	POWER BI, MSSQL
6. Documentation & Portfolio	Project README	GITHUB

Project Timing

The project was executed across six phases covering data ingestion, data cleaning, exploratory data analysis, machine learning, dashboard development, and documentation. For a detailed breakdown of each phase see  [RideMetrics_Timing](#) — Project Timeline

Exploratory Data Analysis (EDA)

A detailed EDA mapping table tracking 12 business questions, chart types, findings and status was maintained separately. All 12 business questions have been answered through visualizations covering earnings optimization, demand patterns, platform comparison and time series analysis. For full findings see  [RideMetrics - EDA Mapping](#).

Tools & Technologies

Programming & Data Processing

- Python — primary programming language for data ingestion, cleaning, feature engineering, EDA and machine learning
- pandas — data manipulation and transformation
- pyodbc — connecting Python to MSSQL for data ingestion
- matplotlib & seaborn — exploratory data visualizations
- scikit-learn — building and evaluating Logistic Regression and Random Forest models
- XGBoost — advanced gradient boosting model used for demand prediction

Database & Data Engineering

- Microsoft SQL Server (MSSQL) — primary database for implementing the Medallion Architecture (Bronze, Silver, Gold layers)
- SQL — querying, transforming and aggregating data across all three layers

Business Intelligence & Visualization

- Power BI — interactive dashboard connected directly to Gold layer tables in MSSQL

Development & Collaboration

- Jupyter Notebook — Python development environment for EDA and machine learning
- GitHub — version control and portfolio hosting
- Notion — project planning and documentation

 [RideMetrics_Timing](#)

 [RideMetrics - EDA Mapping](#)